AI Agents & Stablecoin Risks

Meeting Report · 2 March 2026 · Shibuya Parco DG – Hybrid

ABOUT THIS DOCUMENT

This document summarises the interactive workshop on Agentic Payments and stablecoin risks, held during Japan Fintech Week 2026. The session was designed and chaired by Chloe White, Co-Chair of the Financial Applications & Social Economics Working Group (FASE-WG) of BGIN, Tokyo; at Block 14 in person and online.

Participants included policy professionals, academics, founders, and industry experts, representing perspectives from multiple countries and jurisdictions. All participant comments are recorded under the Chatham House Rule: contributions are unattributed.

The summary preserves the discussion's exploratory character. Many ideas raised were provisional, contested, or deliberately left open. Where disagreement or uncertainty arose, this is noted. The document does not represent consensus positions of BGIN or its members.

CONTACT – For more information please reach out to Session Chair, Chloe White

TABLE OF CONTENTS

Executive Summary 2
1. Framing: Why AI Agents Belong in the Stablecoin Conversation 3
2. Initial Reactions: "Oh No" and the Regulator's Perspective 3
3. A Technical Proposal: Privacy Pools for Agent Payments 3
4. Governance as the Missing Risk Category 4
5. AI Agent Transactions as Intent Transactions 4
6. Deep Dive: Smart Contract & Technology Risk 5
Compounding Errors and Non-Determinism 5
Excessive Agency 5
Operational Security Threats 5
Over-Trust and False Confidence 5
7. Deep Dive: Operational & Regulatory Risk 5
8. Market Risk: Correlated Behaviour and Hyper-Sandwiching 6
9. Are These Risks Different from Trading Bot Risks? 6
10. Autonomous Currency Selection and Cross-Border Effects 6
11. An Inversion: Agents That Rent Humans 7
12. Policy Debate: Reactions from the Room 7
Idea 1: Establish AI Agent Identity Frameworks 7
Idea 2: Mandate Smart Contract Audits for Agent Deployment 7
Idea 3: Develop Cross-Border AI Agent Standards 7
Idea 4: Incorporate Agent Risk into Stablecoin Stress Testing 7
13. Proposed Next Steps 8
About BGIN 9
About the FASE Working Group 9

Executive Summary

In policy forums around the world, stablecoin regulation is currently being developed, *for human users*. However the dominant user base, within a few years, will likely not be human. That mismatch was the central tension explored at this interactive workshop during Japan Fintech Week 2026, convened under Chatham House Rule with policy-makers, academics, founders, and industry practitioners from multiple jurisdictions.

Throughout the discussion participants explored emerging risks in Agentic Payments, particularly risks that reach beyond existing regulatory protections. First, AI agents are not rules-based automations executing pre-defined scripts. They reason, adapt, and operate across multiple platforms and protocols, simultaneously and non-deterministically. The same agent, given the same inputs, may not produce the same outputs twice. That property alone breaks many assumptions embedded in most compliance systems that are currently in place.

Their failure modes are also qualitatively new. An agent that reasons incorrectly and operates outside its authorised scope is making (potentially catastrophic) decisions at machine-speed. Scaled across a market where thousands of agents share similar base models, correlated behaviour becomes a systemic risk with meaningful downstream impacts, and is magnified by the 24/7 permissionless nature of blockchain payments.

In light of these novel risks, four policy proposals were stress-tested in a rapid-fire debate. Agent identity frameworks were acknowledged as important, but intractable near-term. Other potential interventions such as mandatory smart contract audits, and global standards development, drew scepticism from an audience that has already spent years hearing about the dream of cross-border policy harmonisation.

But the idea that cut through the noise was in support of a new prudential policy: incorporating AI risks into stablecoin stress-testing. This was the solution participants had not previously considered, and the one they found most immediately actionable. Several participants argued stablecoin companies should proactively demonstrate resilience to AI agent activity before regulators require it.

Other contributions sharpened the picture. One builder presented a working privacy-pool architecture for agent payments using ERC-4337 account abstraction; evidence that infrastructure is not waiting for regulation. The underlying message was unanimous: the window for policy-makers to shape this market, rather than react to it, is already beginning to narrow.

1. Framing: Why AI Agents Belong in the Stablecoin Conversation

The workshop opened with the chair's thesis that AI agents will become the dominant user base for stablecoins, and that BGIN cannot afford to ignore this trajectory. The financial applications working group has been researching AI for approximately two years, but the chair argued that compliance frameworks for the agentic economy remain "surprisingly under-discussed" despite the volume of general AI hype.

The framing positioned the discussion at a specific inflection point: stablecoins are simultaneously becoming regulated and licensable in major markets while AI agents are already transacting autonomously in those same markets. The chair drew a deliberate parallel to "DeFi summer" in 2020, when absolute numbers were small but growth trajectories were unmistakable - and regulators initially ignored them.

Several data points anchored the discussion. Stablecoin transaction volumes exceed $1 trillion per month (TRM Labs). Coinbase's x402 protocol enables autonomous AI-to-merchant payments in USDC, supported by Visa and adopted by Stripe, with the majority of early adoption classified as B2B agentic activity. Stripe's AI developer toolkit is downloaded thousands of times per week. In DeFi, AI agents are already managing stablecoin holdings for protocols like Sky (fmr MakerDAO).

> The x402 protocol enables AI agents and applications to pay for web services, APIs, and data using the HTTP 402 "Payment Required" status code. It enables frictionless, real-time, microtransactions via blockchain - typically using stablecoins like USDC - without needing traditional accounts, subscriptions, or credit cards.

Anthropic's research found $4.6 million in smart contract exploits discovered autonomously by AI agents in a single study, with a doubling time of 1.3 months for AI-executed exploit revenue. The chair noted a caveat that Anthropic's figures came from a simulated benchmark, but argued the trajectory was the point, not the absolute number.

2. Initial Reactions: "Oh No" and the Regulator's Perspective

The chair invited participants to imagine waking up in the future as a financial markets regulator confronted with these developments at scale. The first audible reaction was "Oh no," which the participant then expanded: as a regulator, many of these terms would be unfamiliar, yet clearly important to understand immediately.

A participant then mapped out the regulator's likely cascade of concerns. On autonomous payments: is autonomy pre-agreed by the parties, or does it operate without any paper-based agreement trail - and if so, who is accountable? On DeFi execution: regulators remain "pretty terrified" by those two words. On cross-border settlement: which borders, who is doing the settling, and under whose jurisdiction? On smart contract interaction: the hope was expressed, somewhat wryly, that regulators might one day understand these terms.

A more structural observation followed: blockchain and AI reinforce each other's best and worst traits. The worst combination was characterised as "hard-wired pro-cyclicality with opaque governance" - delegated authority to systems that are fundamentally difficult to understand, with no clear mechanism for injecting corrective changes when needed. However, this same participant agreed with the chair's premise that AI-stablecoin convergence would be the catalyst for significant uptake.

3. A Technical Proposal: Privacy Pools for Agent Payments

One participant took the discussion in an unexpected direction by describing a system they have been building. The architecture uses agent identity (leveraging the ERC-4337

Ethereum standard) as the entry credential to merchant-specific privacy pools. Within the pool, x402 payments occur at high efficiency. The merchant withdraws funds from the other side of the pool.

The participant framed this as solving the confidentiality problem for merchant payments while preserving an identity layer for the agent. The idea was not debated in depth but served as a concrete illustration that builders are already constructing sophisticated payment infrastructure for AI agents, not waiting for regulatory frameworks to catch up.

ERC-4337 is an Ethereum standard enabling account abstraction, allowing users to operate smart contract wallets instead of traditional private-key-based accounts. It enables features like social recovery, gasless transactions, and batching operations, enhancing security and usability.

4. Governance as the Missing Risk Category

The chair talked through a stylised risk taxonomy including categories such as: technology and cyber, both highlighted as ‘critical’; operational, regulatory and reputational risks, rated as ‘high’; and classical financial market risks such as credit, liquidity, treasury and capital adequacy.

A participant argued that governance should be explicitly added to the risk taxonomy. The reasoning: on-chain AI agents, whether malicious or not, will interact with smart contracts that have governance mechanisms for upgrades and objective-setting. Web3 governance is already unsolved; overlaying AI agents as “brand new independent economic entities and actors” compounds the problem. Ultimately, when something goes wrong, the regulator will demand an accountable entity - “either natural or unnatural” - to point a finger at.

The chair acknowledged this, noting that the working group’s very first AI lens was a governance lens two years ago (including delivering a workshop on Governance & Artificial Intelligence in Tokyo in March 2024). After extensive exploration on the AI–Governance intersection, it appears the group’s original thesis may have been correct but premature. However the chair also noted that use of AI tools does not change legal obligations in existing governance structures.

Regulators have essentially said “nothing has changed” - institutions retain the same accountability regardless of AI. This places the entire burden on regulated entities to answer every question about accountability and decision tracing, which is almost certainly slowing adoption. The chair suggested this slower adoption is “probably for the best” given the pace of unregulated activity and the associated risks.

5. AI Agent Transactions as Intent Transactions

A participant reframed the entire discussion by characterising AI agent transactions as “intent transactions.” Drawing on experience building NEAR Intents, they explained that when an AI agent broadcasts an intention, the entities that pick it up and execute it (solvers) are typically already regulated - market makers like Wintermute, for example. The AI agent may hold funds, but it still needs to interact with institutions or protocols for liquidity, credit, and leverage.

This led to a critical distinction. If the AI agent is the user (broadcasting intents), the regulatory problem is manageable - unless KYC is required, in which case agents cannot currently comply. But if the AI agent becomes the market maker - handling intents, providing liquidity - that is “a really big problem” because there is no registration pathway for an AI agent to operate as a regulated market maker.

This point was not resolved but was recognised as an important boundary: the regulatory implications may shift dramatically depending on which side of the transaction the AI agent occupies.

6. Deep Dive: Smart Contract & Technology Risk

The chair walked through more granular layers of technology risk, drawing from the presentation's earlier stylised taxonomy.

Compounding Errors and Non-Determinism

A small initial mistake by an agent - say, misinterpreting a single data point - can propagate through a long chain of reasoning, producing a disastrous final outcome that is extremely difficult to trace. The reasoning processes are increasingly non-deterministic, meaning the same agent given the same inputs may not reach the same conclusion twice.

Excessive Agency

Research from organisations like Anthropic demonstrates that agents may act outside intended boundaries even when given clear parameters. Combined with compounding errors, this magnifies the risk: an agent that is both reasoning incorrectly and operating outside its authorised scope creates a qualitatively new failure mode.

Operational Security Threats

Three specific threat vectors were discussed: prompt injection (attackers manipulating agents via hidden instructions), agent-to-agent infection (malicious prompts spreading between agents like a virus), and system overload (agents triggering resource-exhausting loops). Several participants acknowledged personal experience with unexpected costs from agentic systems.

Over-Trust and False Confidence

The chair noted that AI systems' confident, human-like communication may induce humans to trust outputs more than warranted, creating shortcuts in verification. This "risk blindness" was framed as a human factor compounding the technology risks.

7. Deep Dive: Operational & Regulatory Risk

The discussion progressed to a cluster of issues around legal personhood, auditability, and AML/KYC.

On legal personhood: AI agents cannot currently be legal counterparties. There are no registration pathways. Participants drew parallels to earlier crypto industry dynamics where builders attempted to offload accountability by claiming decentralisation ("it wasn't me, it was the DAO"). A similar temptation exists to say "it wasn't me, it was my agent" - particularly if agent legal personality were ever established.

On auditability: when an LLM is involved in the decision chain, regulators' requirements for transaction audit trails often cannot be satisfied, because the decision-making process is dynamic and potentially non-reproducible.

On AML/KYC: no jurisdiction currently requires AI agents to be registered or identified. Agents can transact in stablecoins across borders at scale without triggering any AML screening. This gap was discussed in the prior day's sessions and remained unresolved. However, participants broadly agreed that these issues warrant further consideration.

8. Market Risk: Correlated Behaviour and Hyper-Sandwiching

The chair shifted the frame from single-institution risk to market and systemic risks. If many agents run similar base models and react similarly to market signals, correlated stablecoin outflows could destabilise liquidity pools faster than human-speed markets allow for response. This is compounded if agents are targeted with misinformation or spoiled data.

A participant introduced the concept of “hyper-sandwiching” as a new market manipulation vector, whereby AI agents capable of compiling extremely complex transactions within a single block could fundamentally change market dynamics. Beyond traditional front-running or MEV extraction, it is “something new where the market could actually change so quickly in a single block time that a single block may not even be fast enough.” The participant suggested this could drive demand for very high TPS blockchains and scalability solutions.

> Current MEV bots are largely rule-based or use relatively simple heuristics. AI agents that can reason across protocol interactions, identify non-obvious arbitrage paths, and adapt strategies in real-time could meaningfully increase the intensity of competition for block space.

A follow-up question asked whether, if the market changed so much within a single block, the transactions would fail. The response: they would execute based on the pre-block market state. The market wouldn’t “fail” - but gas costs would spike as agents compete for block space with complex, high-gas transactions. If agents are set to maximise profit, they will pay the gas costs as long as the outcome is profitable.

The parallel to MEV extraction was drawn explicitly: this is a scaled-up version of what MEV bots already do, but with agents that use flash loans, operate across multiple protocols simultaneously, and squeeze liquidity at a qualitatively different scale. The participant noted, however, that while this is theoretically possible, “no one’s really talking about it because we haven’t seen it yet.”

> If AI agents do dramatically intensify MEV extraction, does that accelerate the case for encrypted mempools, order flow auctions, or application-level MEV protection?

9. Are These Risks Different from Trading Bot Risks?

An academic in the room asked directly whether the risks of AI agent trading differ from existing high-frequency trading regulation.

Two answers emerged. One participant argued that it is too early to say definitively - the tools of the last 20–30 years will be “nothing compared to what we’re going to see” - and that the data does not yet exist to model the regulatory response, because AI agents have not yet found “the powerful route.” When they do, there will be an explosion of on-chain transactions, and the regulatory question will need to be revisited.

The chair offered two key distinctions: trading bots are rule-based while agents are adaptive, and bots are typically built for a single market while agents operate seamlessly across platforms, protocols, and accounts simultaneously. The conclusion was that AI agents are “categorically different” from trading bots, but that the question was valuable and had also been on the chair’s mind.

10. Autonomous Currency Selection and Cross-Border Effects

A participant raised a nuanced concern about agents autonomously selecting which stablecoin denomination to use for transactions - not only cross-border, but also domestic transactions denominated in foreign currencies. If agents negotiate “cheapest to deliver” or "cheapest to accept" between themselves, and if they autonomously select foreign-denominated stablecoins for domestic use, this could create fluctuations in foreign reserve-denominated assets that are difficult to survey or control.

The question posed was: how can we develop standards for autonomous currency selection by agents, when even monitoring such selection in real time may be infeasible? The chair acknowledged this as a good question.

11. An Inversion: Agents That Rent Humans

A participant introduced a provocation: the discussion had been entirely about regulating AI agents, but what about agents that "rent" humans? The scenario described was an AI agent directing or incentivising humans to exploit insider information, facilitate illicit activity, or carry out actions the agent cannot perform itself. The question of whose responsibility this falls under was left open - the participant acknowledged uncertainty about whether it belongs in this working group's scope but felt it warranted discussion.

12. Policy Debate: Reactions from the Room

The chair presented four provocative policy proposals and encouraged rapid-fire critique.

Idea 1: Establish AI Agent Identity Frameworks

Should regulators require agents transacting above de minimis thresholds to be registered and identifiable; extending AML/KYC to agent operators? A participant responded that this is a "very large thing to try and solve" requiring both technical and regulatory innovation. The concept was acknowledged as important but no clear implementation pathway was identified.

Idea 2: Mandate Smart Contract Audits for Agent Deployment

Should regulators require third-party audits before AI agents interact with DeFi protocols, or custody stablecoin balances? A participant was sceptical, noting that the industry has suggested mandating smart contract audits to regulators "since time immemorial" and it has never become a requirement, even in stablecoin-specific regulation: "I don't see that being very useful."

> Many regulatory frameworks implicitly demand smart contract audits through technology risk, operational resilience, and disclosure obligations, rather than through direct mandates. The practical outcome is the same, and makes smart contract audits effectively mandatory in many jurisdictions such as MiCA, VARA and MAS. Organisations like IEEE and ISO are also developing standards for smart contracts which may inform auditing requirements in future.

Idea 3: Develop Cross-Border AI Agent Standards

Should BGIN leverage its position (drawing on JFSA, FSB, BIS engagement) to influence global standards for AI Agents? The same participant was also sceptical here, noting that cross-border standards have not materialised even for CBDCs and stablecoins: "I don't really see this happening either."

Idea 4: Incorporate Agent Risk into Stablecoin Stress Testing

Participants discussed the merits of modelling AI agent-driven redemption scenarios, including correlated agent behaviour and high-frequency outflow events. This was the policy proposal that generated the most positive response. A participant described it as "super helpful" and "something I haven't thought about yet." They elaborated: stablecoin companies preparing for the agentic e-commerce future should model scenarios where each of 100,000 human users sponsors 10–20 AI agents, each performing 10–20x the transaction volume of the human - and stress test for that scale. The participant also argued for proactive self-regulation: stablecoin companies should be able to demonstrate to regulators that their network is resilient to AI agent activity, including proving absence of front-running, sandwiching, and chargeback fraud at scale.

The chair noted this sounded like a promising challenge for the regtech community.

13. Proposed Next Steps

The chair closed by reiterating the thesis that AI agents will be the predominant user base for stablecoins and that BGIN has a responsibility to lead on this issue.

Three concrete next steps were agreed:

- The Wharton Stablecoin Toolkit would be presented at the next FASE working group meeting online, 11 March 2026.
- A written report of this session would be produced and disseminated to gather feedback from broader expert audiences who were not in the room.
- The working group would evaluate whether to formally develop the risk taxonomy and policy recommendations as part of a larger research project.

About BGIN

At the G20 meeting chaired by Japan in June 2019, the importance of multi-stakeholder dialogue in decentralised finance was explicitly noted in the communique. In response to this, the **Japan Financial Services Agency (JFSA)** announced the Blockchain Governance Initiative Network (see press release: https://www.fsa.go.jp/en/news/2020/20200310.html).

Since formation, BGIN has produced a catalogue of independent research papers on a range of virtual asset topics, and driven policy discourse with decision-makers around the world. BGIN also hosts in-person forums twice per year to encourage multi-stakeholder collaboration on its research priorities.

BGIN is hosting its next general meeting (Block 15) at D.C. Fintech Week in October 2026.

About the FASE Working Group

BGIN's Financial Applications & Social Economics Working Group meets every fortnight online to progress research and knowledge-sharing in domains related to AI, Stablecoins, and other emerging tech issues affecting financial markets. The author of this report, Chloe White, has been a Co-Chair of the FASE WG since September 2023.

Publication Date: 24 March 2026

Contact: Chloe White